PrepAIred: Bounded-Authority Multi-Agent Coordination for Adaptive Technical Interviews

Anonymous Authors

Anonymous Institution

Abstract. Existing technical-interview systems typically solve only part of the interaction problem: static or rubric-based scorers emphasize reproducibility, LLM interviewers emphasize conversational flexibility, and adaptive tutoring systems emphasize personalization. Multi-agent educational systems add role specialization, but the authority and coordination boundaries between decision-making components are often less explicit. We present PrepAIred, a centrally coordinated, bounded-authority multi-agent architecture for adaptive technical interviews. The system combines a deterministic multi-signal technical evaluator, an independent communication-confidence model, a Bloom-aligned question bank, a frozen PPO difficulty policy, stateful Strategy and Coaching agents, typed proposal contracts, evidence-grounded arbitration, and a local Qwen2.5-1.5B-Instruct generator restricted to grounded language production. An exploratory benchmark of 64 constructed answers across eight questions gives the deployed composite $\rho = 0.3812$ against human-reference scores; the CrossEncoder signal gives $\rho = 0.4832$, but its length-controlled partial correlation is 0.2434 (95% CI $[-0.008, 0.507]$), so length-independent reasoning value is not established. In deterministic replay of 160 matched sessions, the coordinated agent layer reduces mean difficulty changes from 15.12 to 12.45 and reversals from 11.74 to 9.87, while adding about 1.04 ms median orchestration overhead per turn. A matched split-versus-unified comparison produces identical decisions on 1,280/1,280 turns. The contribution is therefore a verifiable coordination architecture with bounded authority, not a claim of learning or pedagogical superiority.

Keywords: Multi-Agent Systems · Agent Coordination · Reinforcement Learning · Natural Language Processing · Technical Interview Evaluation · Adaptive Learning

1 Introduction

Technical interviews in computer science require candidates to demonstrate domain knowledge, reasoning, problem solving, and clear communication within a single interaction. Scalable practice systems have therefore evolved along several partially independent directions: static or rubric-based answer assessment, conversational LLM interviewers, adaptive tutoring and testing, and more recently multi-agent educational systems. Each direction addresses a real part of the problem, but combining them creates a systems question that is less well specified:

which component is allowed to make which decision, how are specialized decisions coordinated, and how can the resulting behavior be verified?

Traditional automated assessment methods provide reproducible scores but can be insensitive to conversational context and can over-credit terminology that is not connected to correct reasoning. Transformer-based semantic representations [2] improve coverage, while pairwise scoring models [3] provide stronger contextual similarity signals, but these signals still require explicit rubric constraints to define what counts as a technically sufficient answer. Speech-enabled interview systems add acoustic and temporal information [4, 11], yet delivery quality is often merged into a single overall performance score, making it difficult to distinguish technical correctness from communication state. Adaptive educational systems use item-response theory [8], Bayesian tracing [5], or reinforcement learning [6, 7] to select content, but a live spoken technical interview combines adaptive questioning with open-ended answer evaluation and follow-up dialogue.

Large Language Models (LLMs) changed the interaction layer [9, 10] by making dynamic interview questions and feedback easy to generate. The same flexibility creates a systems risk: when one generative component simultaneously controls scoring, questioning, and feedback, a conversational model can change the effective rubric or difficulty policy through its own language output. Multi-agent LLM frameworks such as AutoGen [12] and MetaGPT [13] show that explicit role decomposition and interaction protocols can organize complex workflows, but they also motivate careful treatment of coordination, authority, and error propagation. Educational multi-agent systems such as MultiTutor [14] and other adaptive tutoring frameworks similarly show the usefulness of specialized roles, but the central research problem in those systems is typically personalized support rather than bounded authority for technical assessment.

PrepAIred addresses this intersection by treating interview preparation as a *coordinated decision system*. The project distinguishes between computational services and explicit agents. A technical evaluator owns numerical technical scoring; a communication-confidence component supplies a separate state signal; a frozen PPO policy proposes difficulty changes; a question selector instantiates valid questions from a structured bank; and a local LLM generates grounded natural-language follow-ups and feedback. On top of this existing pipeline, the specialization configuration studied in this paper introduces two stateful weak agents: a Strategy Agent, which maintains a persistent weakness ledger and proposes difficulty/probe actions, and a Coaching Agent, which maintains coaching history and proposes feedback depth and remediation emphasis. Neither agent can write session state directly. A deterministic coordinator validates their typed proposals, checks evidence and version freshness, arbitrates conflicts against the frozen PPO action and guardrails, and commits the result.

This architecture is intentionally *centrally coordinated*, not decentralized multi-agent reinforcement learning. The design goal is not to maximize agent autonomy. It is to make specialization explicit while preserving least-privilege authority, auditable state ownership, deterministic safety envelopes, and reproducible behavior.

This distinction is important for systems where incorrect scores or uncontrolled difficulty changes can directly affect a candidate's learning experience.

The work makes the following contributions.

1. **A bounded-authority multi-agent architecture.** We formalize a centrally coordinated interview system in which specialized agents maintain private state, return typed proposals, and communicate through explicit contracts while a single coordinator remains the sole system of record.
2. **A coordination and arbitration protocol.** Agent proposals pass schema, authority, version, and evidence checks before arbitration. Threshold gating and deterministic guardrails bound their effects; invalid or timed-out proposals are dropped.
3. **A cognitively structured adaptive question bank.** Each technical question carries difficulty and Bloom-taxonomy metadata. The operational mapping used by the project spans Understand, Apply, Analyze, Evaluate, and Create bands, while the adaptive policy separately operates on five bounded numeric difficulty tiers.
4. **A technically grounded evaluation pipeline.** PrepAIred separates technical scoring, communication state, adaptive strategy, question selection, and natural-language generation. The technical evaluator uses semantic similarity, rubric concept coverage, and a CrossEncoder-derived signal, with mandatory-concept and mistake controls.
5. **Empirical evaluation of coordination and specialization.** Deterministic replay evaluates baseline, coordinated, silent-agent, arbitration-ablated, peer-messaging-ablated, PPO-ablated, and unified-controller conditions across 160 matched sessions derived from eight underlying replay clusters and 20 question-order permutations.
6. **Transparent characterization of evaluator validity.** An exploratory benchmark of 64 constructed answers shows moderate association between the deployed evaluator and human-reference scores, but also demonstrates a substantial answer-length confound. We therefore do not claim length-independent reasoning superiority or candidate learning gains.

Research questions. The study addresses four questions: (RQ1) Does stateful agent coordination alter interview difficulty trajectories relative to the frozen single-policy baseline? (RQ2) Does decomposing Strategy and Coaching into separate agents produce different decisions than an equivalent unified controller with matched information? (RQ3) Which coordination mechanisms—stateful memory, peer messaging, confidence-threshold arbitration, and PPO interaction—account for observable behavior? (RQ4) How closely do the frozen technical-evaluator scores agree with the available human-reference benchmark, and do they add evidence beyond answer length and simple lexical baselines?

2 Related Work

2.1 Automated Technical Answer Assessment

Automated short-answer and technical-answer scoring has progressed from keyword matching to contextual neural representations. Sentence-BERT produces efficient sentence embeddings for semantic similarity [2], while CrossEncoder architectures score pairs jointly and can represent richer contextual interactions. These methods help with coverage but do not, by themselves, define the authority of a score inside a multi-turn system. PrepAIred therefore treats the evaluator as a protected computational service whose score is an input to other decision components rather than an instruction for a generative model.

2.2 Adaptive Learning, Reinforcement Learning, and Question Difficulty

Adaptive testing and tutoring have long selected content according to an estimate of learner state. Item-response approaches optimize question difficulty from response patterns [8], while RL formulations optimize action selection over trajectories. PrepAIred uses PPO [6] only for the difficulty proposal and keeps hard bounds and deterministic guardrails outside the learned policy. The question bank adds an orthogonal cognitive-structure layer based on the revised Bloom taxonomy, whose levels are commonly used to organize learning objectives and question complexity. Published educational assessment work has also used Bloom categories as a structured basis for question complexity and difficulty.

2.3 Multi-Agent Systems and Coordination

Classic agent theory emphasizes autonomy, reactivity, pro-activeness, and social ability [1]. Modern LLM-based frameworks such as AutoGen enable flexible conversation among multiple agents, whereas MetaGPT encodes structured operating procedures and role specialization into an assembly-line workflow. These systems demonstrate the usefulness of decomposition but also expose coordination and cascading-error risks [15]. Recent educational multi-agent work similarly uses specialized tutoring roles to provide adaptive support [14]. PrepAIred differs in emphasizing *authority separation*: specialized agents propose actions but cannot directly change scores, difficulty, or session state.

2.4 Agent Safety, Governance, and Verification

The principle of least privilege [16] provides a useful abstraction for agent authority: components should receive only the permissions required for their role. Safe RL shielding [17] further motivates keeping hard constraints external to a learned policy. PrepAIred combines these principles with typed state contracts, version checks, evidence grounding, timeout isolation, and a coordinator that is the sole state writer.

Table 1. Positioning of PrepAIred against representative approaches.

Theme	**Representative approaches**	**Typical emphasis**	**PrepAIred focus**
Answer assessment	Sentence embeddings, CrossEncoders	Semantic or pairwise scoring, often static	Protected multi-signal evaluator inside a live adaptive loop
Adaptive sequencing	ITS, item response, PPO	Learner-state estimation and next-item selection	PPO difficulty proposal with external guardrails
LLM interviewers	Generative interviewers / LLM judges	Flexible dialogue and feedback	LLM restricted to grounded wording and feedback
Educational MAS	MultiTutor, adaptive tutoring MAS	Specialized pedagogical support	Persistent Strategy/Coaching roles with bounded authority
Generic MAS	AutoGen [12], MetaGPT [13]	Agent conversation or structured role workflows	Typed proposals, state isolation, deterministic arbitration

2.5 Positioning and Gap

The closest bodies of work solve complementary portions of the problem: automated answer scoring, adaptive tutoring, LLM interviewing, or multi-agent personalized support. The contribution here is the integration of these roles into a *bounded-authority, centrally coordinated* architecture for live technical-interview preparation, together with empirical tests of trajectory effects, specialization limits, coordination mechanisms, and evaluator confounds. A focused literature comparison did not identify a directly comparable system combining all of these elements in one implemented setting; this is stated as a scoped search result, not as an absolute novelty claim.

3 Problem Formulation and System Scope

PrepAIred models a session as a sequence of primary interview turns with optional subordinate follow-up probes. The system must simultaneously solve four coupled problems: (i) infer technical evidence from a free-form answer, (ii) estimate a communication-confidence state, (iii) choose the next difficulty subject to hard bounds and session constraints, and (iv) provide grounded natural-language coaching. A monolithic controller could solve all four tasks procedurally, but this would obscure authority boundaries and make it difficult to test specialization independently.

We therefore separate *observation*, *proposal*, *arbitration*, and *commit*. The coordinator constructs a validated turn observation from the current session state.

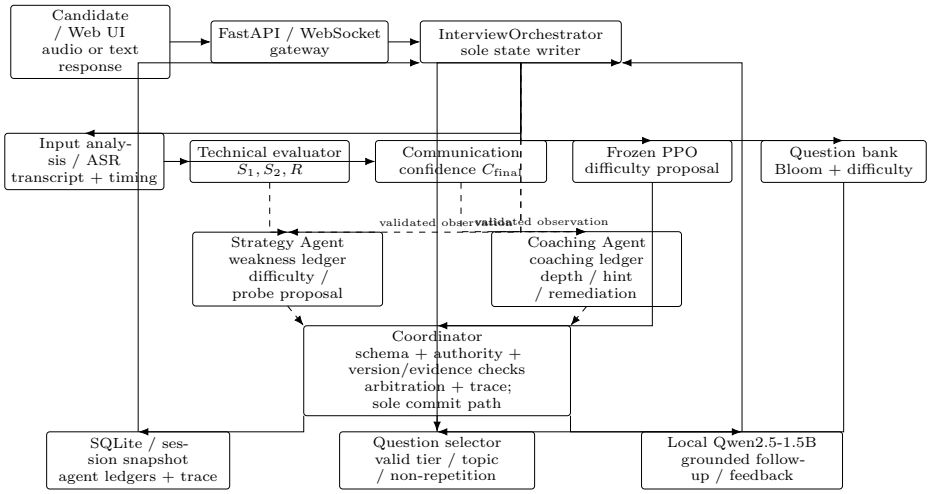


Fig. 1. Bounded-authority multi-agent architecture of PrepAIred. The evaluator, communication model, frozen PPO policy, Bloom-aware question bank, and local LLM are specialized system components; the evaluated agent layer adds stateful Strategy and Coaching roles. The coordinator is the sole state writer and applies schema, authority, version, and evidence checks before arbitration and commit. Dashed links denote validated observations or bounded agent proposals.

Specialized components consume that immutable observation. Agent proposals are returned as typed objects and checked before any proposal can affect session progression. The coordinator applies deterministic precedence: infrastructure and hard constraints first, guardrails next, a validated high-confidence Strategy proposal next, and the frozen PPO action or heuristic fallback last. Feedback and follow-up generation operate only on evidence already produced by the evaluator and coordinator.

This scope deliberately excludes claims of decentralized MARL, emergent agent negotiation, or human learning gains. The evaluated multi-agent layer is a centrally orchestrated specialization architecture around a frozen PPO decision component.

4 PrepAIred Architecture

The system is a modular pipeline with explicit boundaries between candidate interaction, observation generation, evaluation, adaptive decision-making, question retrieval, agent coordination, and natural-language generation. The main runtime architecture and its session control flow are shown in Fig. 1.

Table 2. Role and authority matrix in the specialization configuration.

Role	Private state / knowledge	May propose	May not do
Input analysis	Transcript, timing, response features	Structured observation fields	Change score, difficulty, or session state
Technical evaluator	Rubric, semantic/concept/misconception evidence	Technical score and evidence	Invoke agent actions or alter session state
Communication model	Acoustic/linguistic state	Confidence/hesitation state	Alter technical score through the acoustic channel
Strategy Agent	Weakness ledger, probe history, score trend	Difficulty action; probe target/type; strategy mode	Write score, set tier directly, terminate session, override hard constraints
Coaching Agent	Coaching ledger, prior hints/examples, repeated misconceptions	Feedback depth, hints/examples, emphasis, remediation tag, tone	Change evaluator results, request difficulty changes, terminate session
Frozen PPO	Learned policy over 6-D state	Easier/Same/Harder action	Bypass guardrails or write state
Coordinator	Committed session state and versions	Arbitrated plan/action; trace records	Delegate sole state ownership
Qwen2.5-1.5B	Grounded prompt context and evaluator evidence	Follow-up wording; qualitative feedback	Determine score, grade, or difficulty

4.1 Roles and Authority

The project distinguishes computational services from explicit agents. This avoids treating every software module as an agent simply because it participates in the workflow. The two explicit weak agents studied in the specialization configuration are stateful, proactive, and communicative, but their authority is deliberately bounded.

4.2 Typed Communication and State Ownership

Each agent receives an immutable, validated turn observation containing the technical score, evidence lists, current difficulty, score history, progress, and follow-up counters. It returns one typed proposal. Unknown fields, non-finite values, out-of-range values, stale turn versions, and ungrounded concepts are rejected. The coordinator records the session, turn, sender, recipient, state version, provenance, and decision outcome. A proposal that times out or raises an exception is dropped; the default policy remains responsible for the turn.

4.3 Question Bank and Bloom-Aligned Difficulty Levelization

The question bank stores structured metadata for each question, including question identifier, topic, numeric difficulty, Bloom level, question type, and time

Table 3. Operational Bloom-aligned bands used in the question bank.

Bloom level	Operational normalized band	Role in question design
Understand	0.20–0.40	Explain or interpret a known concept
Apply	0.40–0.60	Apply a concept to a concrete task
Analyze	0.60–0.80	Decompose, compare, trace, or diagnose
Evaluate	0.80–0.90	Judge alternatives or trade-offs using criteria
Create	0.90–1.00	Synthesize or design a solution

limit. Bloom taxonomy [19, 20] provides a cognitive-structure layer above the numeric difficulty tiers. The project uses the following operational mapping for its question metadata:

The Bloom band is a cognitive-design attribute; the adaptive controller separately operates on a bounded numeric difficulty tier $D \in \{1, 2, 3, 4, 5\}$. In the current workflow, the first two primary questions establish a deterministic Easy-pool baseline, after which the controller can move at most one difficulty tier per adaptive decision. Thus Bloom level and numeric difficulty are complementary rather than interchangeable: Bloom metadata constrains the intended cognitive demand of an item, while the RL/strategy layer selects among valid difficulty tiers. The mapping is an engineering specification and has not been presented as a learned calibration from human difficulty ratings.

4.4 Technical Evaluator

The evaluator provides the quantitative evidence used by the coordinator, Strategy Agent, and follow-up gate. It is deliberately independent from the local LLM.

4.5 Communication Confidence

Speech-derived communication state is computed separately from technical correctness and is used as an adaptation signal rather than as a direct technical penalty. The model combines acoustic, linguistic, pause, rate, rhythm, and ASR confidence features. A narrow text-level hedge penalty may enter technical scoring only when content is already weak; acoustic confidence itself is not a parameter of the technical evaluator.

4.6 Adaptive Policy and Question Selection

The PPO policy consumes a normalized six-dimensional state containing current performance, rolling performance, communication confidence, hesitation, the implementation's normalized session-progress component, and current difficulty.

The action space is {Easier, Same, Harder}; transitions are clipped to the valid difficulty range. The project records describe a reward combining rule-oracle alignment, outcome delta, multimodal shaping, and stability terms. Because the frozen training evidence is based on a single simulated persona and does not establish superiority over the rule fallback, this paper treats PPO as one decision input rather than claiming a learned performance advantage.

4.7 Grounded Generation and Follow-up

The follow-up decision is deterministic and evidence gated. A clean, high-scoring answer without a missing concept, misconception, or gap exits the follow-up loop; otherwise, structured evidence can trigger up to two bounded follow-up attempts. The local Qwen2.5-1.5B-Instruct model [10] receives the original question, candidate response, and structured missing-concept/misconception evidence. It generates language only. It cannot alter the stored score, grade, or difficulty action. Candidate-facing feedback suppresses internal numerical values.

5 Technical Evaluator

The technical evaluator compares a candidate transcript against a question-specific rubric containing semantic reference statements, required concepts, misconception vectors, advanced-insight vectors, and mandatory-concept flags. The design is intentionally multi-signal so that no single lexical or semantic measure determines the score.

5.1 Semantic Similarity

Let T_c be segmented into sentences $\{s_1, \ldots, s_n\}$ and let R_q^{semantic} contain reference statements v_j. With L2-normalized embeddings from `all-MiniLM-L6-v2`,

$$S_1 = \frac{1}{|R_q^{\text{semantic}}|} \sum_j \max_i \cos(e_i, v_j). \tag{1}$$

5.2 Concept Coverage and Reasoning Gate

For each required concept c_k, coverage is

$$S_2 = \frac{1}{|R_q^{\text{concept}}|} \sum_k \mathbf{1}\Big(\max_i \cos(e_i, c_k) > 0.30\Big). \tag{2}$$

The stored reasoning signal is used as a gate: $S_2^{\text{effective}} = S_2$ when $R > 0.30$ and $0.60 S_2$ otherwise. This mechanism is designed to reduce inflated coverage from terminology overlap when reasoning evidence is weak; it is not described as proving elimination of keyword stuffing.

Table 4. Technical evaluator parameters in the tested configuration.

Parameter	Value	Purpose
S_1 / S_2 / R weights	0.15 / 0.35 / 0.50	Semantic, concept, and reasoning contributions
Concept threshold	0.30	Required-concept match cutoff
Reasoning floor	0.30	Activates concept-coverage dampening
Dampening factor	0.60	Reduces effective coverage when reasoning is weak
Bonus increment / threshold	0.03 / 0.50	Advanced-insight credit
Mistake increment / threshold	0.07 / 0.55	Misconception penalty
Penalty cap	0.30	Maximum accumulated penalty
Mandatory cap	0.60	Ceiling when a required concept is missing

5.3 CrossEncoder Signal

A partially fine-tuned derivative of `cross-encoder/ms-marco-MiniLM-L-6-v2` scores the candidate/reference pair. The raw model output is treated as a regression score and mapped by

$$R = \mathrm{clamp}\left(\frac{r_{\mathrm{raw}} - 0.20}{0.70}, 0, 1\right). \tag{3}$$

The model's training provenance is incomplete in the project records. The paper therefore reports its empirical behavior without assigning the observed signal a stronger semantic interpretation than supported by the data.

5.4 Score Aggregation and Constraints

The deployed weighting is

$$S_{\mathrm{base}} = 0.15S_1 + 0.35S_2^{\mathrm{effective}} + 0.50R. \tag{4}$$

Advanced-insight matches contribute +0.03 each at a 0.50 cosine threshold; misconception matches contribute −0.07 each at a 0.55 threshold, with accumulated mistake penalties capped at 0.30. If a mandatory concept is required and missing, the final score is capped at 0.60. The project also applies a narrow capped hedge penalty when lexical uncertainty cues co-occur with already-weak content.

6 Communication Confidence Model

The communication model is independent of technical correctness through the acoustic channel. It uses jitter, shimmer, HNR, pitch variation, pause timing,

speaking rate, rhythm, ASR confidence, and a linguistic-confidence signal. The principal equations are:

$$C_{\text{jitter}} = \text{clamp}(1 - j/0.03, 0, 1), \quad C_{\text{shimmer}} = \text{clamp}(1 - s/0.12, 0, 1), \quad (5)$$

$$C_{\text{HNR}} = \text{clamp}(h/28, 0, 1), \quad (6)$$

$$V_Q = 0.32C_{\text{jitter}} + 0.30C_{\text{shimmer}} + 0.28C_{\text{HNR}} + 0.10C_{\text{pitch var}}, \quad (7)$$

$$C_{\text{final}} = 0.28V_Q + 0.27L + 0.22F_{\text{fluency}} + 0.13F_{\text{rate}} + 0.10F_{\text{rhythm}}. \quad (8)$$

If ASR confidence is below 0.45, the implementation shifts part of the score toward V_Q. Communication confidence is used as a state signal for adaptation and is not exposed as a numeric candidate-facing score. The model is not presented as a clinically or psychologically validated measure of confidence.

7 Adaptive Difficulty and Bloom-Aware Question Selection

7.1 State, Action, and Difficulty

The current implementation uses a six-dimensional state consisting of current performance, rolling performance, communication confidence, hesitation, normalized session progress, and current difficulty. The state is normalized to $[0, 1]$. The PPO action set is

$$a_t \in \{\text{Easier}, \text{Same}, \text{Harder}\}, \quad (9)$$

with the corresponding difficulty update bounded to $D \in \{1, 2, 3, 4, 5\}$. The first two primary questions form the baseline phase; adaptive selection begins from Question 3. Follow-ups are subordinate turns and do not increase the primary-question counter.

7.2 Bloom Taxonomy as a Cognitive Difficulty Layer

Bloom's taxonomy [19, 20] provides a complementary interpretation of difficulty. The classic taxonomy and its revision organize learning objectives by increasing cognitive complexity. In PrepAIred, the project uses a five-band operational mapping from Understand through Create (Table 3) while retaining a separate five-tier numeric difficulty index for the adaptive policy. Thus the same technical topic can be represented at different cognitive levels: for example, explaining a known concept can occupy an Understand band, applying it to a code fragment an Apply band, diagnosing a failure an Analyze band, comparing design trade-offs an Evaluate band, and designing a new system an Create band.

This separation is useful for multi-agent coordination: the Question Selector owns item retrieval, the Strategy/PPO layer owns difficulty movement, and the Bloom metadata remains a structural property of the candidate item rather than a quantity that an LLM is allowed to rewrite during the session.

7.3 Reward Scope and Evidence

The frozen PPO checkpoints were trained using a hybrid reward based on alignment to a hand-coded oracle, outcome delta, multimodal shaping, and stability. Because the training configuration uses a single simulated persona and unseeded candidate noise, the paper does not infer cross-candidate generalization from those runs. The coordination study instead evaluates the frozen policy as one proposal source inside a deterministic replay environment.

8 Agent Coordination, Arbitration, and Verification

8.1 Proposal Lifecycle

For each turn, the coordinator creates a versioned observation. Strategy and Coaching receive immutable views and return typed proposals. Each proposal passes four sequential checks: (1) strict schema validation with unknown fields rejected and NaN/out-of-range values refused; (2) forbidden-key scanning to prevent direct mutation of protected state; (3) optimistic concurrency verification against the committed turn version; and (4) grounding checks requiring proposed concepts to exist in evaluator evidence. An exception or timeout (150 ms in the tested configuration) drops the proposal and leaves the default policy active.

8.2 Arbitration Precedence

The deterministic precedence order is: (1) infrastructure constraints such as evaluator failure or difficulty bounds; (2) the guardrail layer, including rules that may prevent or force an action; (3) a validated Strategy proposal whose confidence exceeds $\tau = 0.6$ and which cites supporting evidence; and (4) the PPO action, or the labelled heuristic fallback if PPO is unavailable. A proposal that cannot take effect at the current boundary is recorded as modified. The coordinator remains the sole state writer.

8.3 Peer Messaging

The Strategy Agent may pass a typed peer message to Coaching. The message expresses pedagogical intent (for example, remediation or stretch mode) but does not expose private strategy state. In the tested configuration, this peer message changes coaching presentation rather than feeding back into Strategy's difficulty decision. This separation lets the paper test whether communication itself changes system trajectories rather than assuming that it does.

9 Experimental Methodology

9.1 Study A: Evaluator Evidence

Study A uses the 64-answer exploratory human-reference benchmark already present in the project records. The eight question clusters contain constructed

technical answers; 54 answers use the mean of three recorded ratings, and 10 use recorded adjudicated values for cases where rater disagreement exceeded the working criterion. Rater provenance, formal independence, and institutional ethics status are incompletely documented; the benchmark is therefore not described as an expert-panel or confirmatory evaluation.

The three-rater inter-rater reliability was ICC(2,1)=0.953, but the benchmark is strongly separated by construction category. The paper therefore emphasizes cluster-aware inference and treats the results as evidence about this benchmark rather than as population validity.

For load-bearing statistics, the project uses a two-level bootstrap with seed 42 and $B = 10{,}000$: resample the eight questions with replacement and then resample the eight answers associated with each selected question. Paired signal contrasts use the same bootstrap draw for all compared signals.

9.2 Study B: Multi-Agent Coordination Replay

The coordination study uses the real InterviewOrchestrator in deterministic replay mode with evaluator outputs fixed. The replay corpus contains 64 technical-interview replay instances covering arrays, dynamic programming, trees/graphs, system design, concurrency, databases, networking, and object-oriented design. Each underlying replay cluster is combined with 20 seeded question-order permutations (seeds 42–61), yielding 160 sessions per condition. These 160 sessions are not independent; the eight underlying replay clusters are treated as the independent structural units in cluster bootstrap analyses using the clustered-data methodology of Field and Welsh [18]. No live generative LLM is invoked during replay; feedback and follow-up text are produced by deterministic templates so that measured trajectory differences isolate coordination logic.

9.3 Baseline and Ablation Conditions

The primary baseline B0 is the original workflow without agents: frozen PPO (or fallback heuristic) plus guardrails and static follow-up thresholds. M1 is the full coordinated Strategy+Coaching system. M2 is the silent-agent condition. M3 removes Coaching. M4 removes arbitration checks. M5 removes peer messaging. M7 disables PPO. Additional conditions compare canonical and deployed guardrail variants. The specialization question is evaluated by comparing the two-agent architecture against a unified controller that receives the same observations and uses equivalent decision thresholds.

9.4 Metrics

We report session completion and authority invariants; difficulty-level changes and reversals per session; action agreement with PPO; proactive follow-up turns; end-to-end coordination latency; agent calls per session; decision agreement between split and unified architectures; and sensitivity to arbitration threshold

Table 5. Per-signal association with human-reference scores ($n = 64$, eight question clusters).

Signal	ρ	95% CI
Semantic similarity S_1	0.2070	$[-0.0231, 0.5027]$
Concept coverage S_2^{eff}	0.3006	$[-0.0057, 0.6523]$
CrossEncoder R	0.4832	$[0.2453, 0.6956]$
Deployed composite	0.3812	$[0.1529, 0.6473]$

τ. All confidence intervals over replay effects are cluster-bootstrap percentile intervals across the eight replay clusters. They are descriptive rather than formal hypothesis tests.

10 Results

10.1 Evaluator Evidence

The frozen CrossEncoder-only score has Spearman $\rho = 0.4832$ (case-bootstrap 95% CI $[0.250, 0.676]$), while the deployed composite reaches $\rho = 0.3812$. The per-signal associations are summarized in Table 5.

The paired question-aware contrasts are inconclusive at this sample size: $R - S_1 = 0.2761$ with CI $[-0.0747, 0.5745]$, $R - S_2^{\text{eff}} = 0.1826$ with CI $[-0.1632, 0.5353]$, $S_1 - S_2^{\text{eff}} = -0.0936$ with CI $[-0.2938, 0.1654]$, and composite$-R = -0.1020$ with CI $[-0.2853, 0.1171]$. None excludes zero.

A word-count-only baseline reaches $\rho = 0.4897$ (95% CI $[0.219, 0.708]$), while BM25, TF-IDF, and reference-token overlap reach 0.3812, 0.245, and 0.4301, respectively. The raw CrossEncoder score is strongly associated with answer length ($\rho = 0.6594$, 95% CI $[0.457, 0.831]$). Controlling for length reduces the CrossEncoder's human-reference partial correlation to 0.2434 (95% CI $[-0.008, 0.507]$). The paper therefore does not claim a demonstrated length-independent reasoning contribution. The project reports a paired fine-tuned-vs-upstream Spearman difference of 0.337 (95% CI $[0.042, 0.606]$), showing that the fine-tuning changed the measured association, but incomplete training provenance prevents stronger attribution.

10.2 Coordination Dynamics

All 160 sessions in each tested condition completed, committed scores matched recorded evaluator outputs, and difficulty remained within $[1, 5]$ with one-level transition bounds. In the silent-agent condition M2, trajectories reproduced the baseline exactly, providing a useful transparency check.

Enabling the full coordinated system M1 changed the difficulty path in all 160 replay sessions. Mean level changes decreased from 15.119 to 12.45, a difference of -2.67 with cluster-bootstrap 95% CI $[-4.50, -0.24]$. Mean directional reversals

Table 6. Primary coordination results from deterministic replay (160 sessions per condition).

Metric	B0	M1	Observed change / interpretation
Level changes/session	15.119	12.450	-2.67; cluster CI $[-4.50, -0.24]$
Reversals/session	11.744	9.869	-1.88; cluster CI $[-3.12, -0.27]$
Proactive probes/session	0	+0.59	Cluster CI $[0.00, 1.76]$; driven by one cluster
Median turn latency	0.546 ms	1.589 ms	+1.04 ms coordination overhead
Agent calls/session	–	47.7	Coordination cost in the tested configuration
Silent-agent trajectory	baseline	identical	M2 demonstrates transparent infrastructure

decreased from 11.744 to 9.869, a difference of -1.88 with CI $[-3.12, -0.27]$. Proactive probing added +0.59 turns per session, with cluster CI $[0.00, 1.76]$; however, all of that additional probing came from one of the eight replay clusters, so the magnitude should not be generalized.

The multi-agent layer increased median end-to-end turn latency by about 1.04 ms (95% CI $[1.00, 1.08]$ ms) and executed about 47.7 agent calls per session. This is a measurable coordination cost, not a claim of practical infeasibility.

10.3 Limits of Agent Specialization

The two-agent specialization and the matched unified controller produced identical decision outputs on all 1,280 replay turns. Likewise, removing Coaching did not alter the Strategy-driven difficulty trajectory because Coaching has no authority or information pathway into difficulty arbitration. These findings do not invalidate the multi-agent architecture; they show that, under deterministic rules and matched information, decomposition alone did not introduce a different decision policy. The measurable value is therefore architectural—state isolation, role separation, auditability, and bounded authority—rather than a demonstrated superiority of one controller over another.

10.4 Coordination Mechanism Sensitivity

The arbitration threshold was swept over $\tau \in \{0.0, 0.5, 0.6, 0.7, 0.8, 0.9\}$ while holding evidence checks and guardrails active. At $\tau = 0.0$, mean level changes collapsed to 2.225 and PPO agreement to 0.014, reproducing the behavior of removing arbitration. For $\tau = 0.5$–0.8, the observed trajectory remained at 12.45 changes and 9.869 reversals; at $\tau = 0.9$, the Strategy Agent was silenced and the trajectory returned to the baseline. The threshold is therefore a critical mechanism in the tested coordination envelope.

Removing PPO changed difficulty trajectories in only 4 of 160 sessions (2.5%) in the registered replay condition, revealing strong behavioral alignment between the Strategy heuristic and the deterministic fallback. This negative result is important: the paper does not claim that the learned PPO policy provides a large marginal contribution under these frozen replay conditions.

10.5 Peer Messaging and Unified Control

Removing the Strategy-to-Coaching peer message produced difficulty trajectories identical to M1 across all 160 sessions. The peer message affected coaching presentation—depth, remediation labels, and emphasis—but did not feed back into Strategy's difficulty decision. In the unified-controller comparison, split and unified architectures agreed on all 1,280 turns. Thus the experiment distinguishes *coordination infrastructure* from *decision diversity*: the former is present and measurable, while the latter is not demonstrated under the tested deterministic configuration.

11 Discussion

The results support a careful interpretation of multi-agent design for adaptive technical-interview systems. The Strategy Agent's persistent weakness ledger changes the trajectory of difficulty transitions, smoothing repeated up/down oscillations relative to the frozen baseline. This is an algorithmic behavior change, but it is not a pedagogical outcome: deterministic replay cannot establish that smoother trajectories improve learning, confidence, retention, or interview performance.

The more important systems result is the explicit separation of authority. The evaluator owns numerical technical evidence; the PPO policy proposes difficulty; Strategy can propose changes but cannot write the difficulty state; Coaching can modify feedback presentation but not technical outcomes; and the LLM can generate language but cannot choose scores or difficulty. This makes conflicts observable and gives the coordinator a deterministic place to enforce hard constraints. The same property also explains why the architecture remains stable when agents are silent or invalid: proposals can be dropped without changing ownership of the session state.

The specialization comparison is deliberately negative. Splitting Strategy and Coaching did not change the matched decision trajectory relative to a unified controller with equivalent information. A reviewer might therefore ask why the split is necessary. The evidence supports an architectural answer, not a performance answer: the split isolates domain-specific state ledgers, makes responsibility explicit, enables independent audit records, and creates a surface on which coordination mechanisms can be experimentally isolated. At the same time, the split introduces measurable orchestration cost. The paper consequently avoids the stronger claim that multi-agent decomposition is intrinsically better.

The Bloom layer plays a different role. It does not make the agents autonomous; it gives the Question Selector a structured cognitive attribute with which to constrain candidate items. This lets adaptive difficulty operate over both a numerical tier and a stable cognitive-design attribute. Existing research has used Bloom taxonomy to classify or calibrate question difficulty; PrepAIred uses the taxonomy as metadata in a live adaptive interview workflow rather than as a standalone classifier.

The evaluator results reinforce the same design philosophy. The technical score has transparent, reproducible structure, but the exploratory human benchmark is too small and confounded to support a strong claim about reasoning quality. The CrossEncoder component correlates with the human-reference score, yet the word-count baseline is comparable, and controlling for length removes much of the apparent association. The system therefore benefits from a protected deterministic evaluator architecture even while the scientific validation of one component remains incomplete. In a multi-agent context, this separation is important: the generative agent can be useful without being trusted with numerical authority.

12 Limitations and Threats to Validity

The main limitations are methodological and architectural.

- **Deterministic replay.** The 160-session coordination study reuses eight underlying replay clusters across question-order permutations. Cluster bootstrap reduces artificial precision, but eight clusters remain a small structural sample, and replay cannot capture dynamic responses from real candidates.
- **No live LLM in coordination experiments.** The replay study stubs the LLM gateway and uses deterministic template generation. Therefore, the measured coordination effects do not establish live LLM quality, hallucination rates, or human-facing coaching quality.
- **No learning-outcome study.** The work does not establish candidate learning gains, retention, interview placement success, or user satisfaction.
- **Agent specialization scope.** The Strategy and Coaching agents are rule- and ledger-based, not separately trained LLM agents. The findings therefore concern bounded weak agency and deterministic coordination, not open-ended emergent MAS behavior.
- **Evaluator validation.** The 64-answer benchmark is exploratory, internally constructed/recorded, and has incomplete rater-provenance and ethics documentation. Its eight-question cluster structure limits power for fine-grained comparisons.
- **CrossEncoder provenance.** The training data, exact split, and complete training code for the derivative CrossEncoder are not recoverable from the project record, so the paper reports measured behavior without claiming a specific learned reasoning property.
- **PPO evidence.** The frozen PPO evidence derives from a single simulated persona configuration with unseeded candidate noise; cross-persona generalization and superiority over simpler policies are not established.

- **Implementation snapshot.** Some regression fixes and feature-flagged coordination code existed in a reviewed development snapshot whose merge status must be synchronized with the exact submission revision. The paper should therefore be submitted from a source snapshot that reproduces the reported configuration.
- **Confidence model.** Speech-confidence features are sensitive to acoustic conditions and are not presented as a validated psychological measure of confidence.

13 Ethics, Provenance, and Data Availability

No new human-subjects experiment was conducted specifically for the coordination study. The evaluator benchmark was prepared and rated within the project workflow, and its records do not provide a separately documented institutional ethics determination or complete rater provenance. The replay corpus is used as offline software-evaluation input and is stripped of names and direct identifiers in the retained payloads. The paper therefore avoids claims of independent expert evaluation or population-level validity.

The reported coordination study is reproducible only to the extent that the project source snapshot, deterministic replay inputs, and configuration flags are preserved. The authors should release, or make available under appropriate conditions, the de-identified replay metadata and experiment scripts needed to regenerate the reported tables before camera-ready submission if institutional and project governance permits.

PrepAIred uses React 18, FastAPI/WebSockets, PyTorch-based evaluation components, SentenceTransformers, FAISS-CPU, Stable-Baselines3 PPO, and local GGUF/llama.cpp inference with Qwen2.5-1.5B-Instruct. The question bank stores rubric and Bloom metadata, while the coordinator persists primary-turn state and agent ledgers. The technical evaluator and coordination study use fixed model/configuration snapshots for the reported experiments. The code-submission interview path is separate from verbal scoring: code answers are executed against fixed tests in an isolated sandbox rather than scored by the verbal evaluator.

14 Conclusion

PrepAIred demonstrates how adaptive technical-interview practice can be organized as a centrally coordinated, bounded-authority multi-agent system rather than as an LLM-only interviewer. The architecture combines a deterministic technical evaluator, a separate communication-confidence signal, Bloom-aligned question metadata, a frozen PPO difficulty proposal, stateful Strategy and Coaching agents, typed coordination contracts, and a local generative model restricted to grounded language production.

The empirical findings are deliberately mixed. The coordinated agent layer changes difficulty trajectories and reduces oscillation in deterministic replay,

but it adds orchestration overhead, and the matched split-versus-unified comparison produces identical decisions. The strongest interpretation is therefore architectural: explicit specialization provides isolated state ownership, proposal contracts, auditable arbitration, and fault containment, while the experiments do not establish that agent decomposition intrinsically improves decision quality. The evaluator benchmark reaches moderate agreement with human-reference scores but also shows that answer length and simple lexical signals explain much of the observed association, so no strong length-independent reasoning claim is made.

Future work should use new human-authored, preregistered data to test the evaluator beyond length confounds; compare the specialization layer against stronger learned unified and multi-agent baselines; evaluate live LLM-backed agents with human participants; and investigate whether bounded authority remains beneficial under larger agent populations, richer peer communication, and decentralized coordination. These extensions preserve the central design principle of the current system: agent autonomy should grow only when the evidence and authority boundaries required for safe coordination grow with it.

References

1. M. Wooldridge, N.R. Jennings: Intelligent Agents: Theory and Practice. The Knowledge Engineering Review 10(2), 115–152 (1995). doi:10.1017/S0269888900008122
2. N. Reimers, I. Gurevych: Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In: Proceedings of EMNLP-IJCNLP, pp. 3982–3992 (2019)
3. P. Bajaj et al.: MS MARCO: A Human Generated Machine Reading Comprehension Dataset. arXiv:1611.09268 (2016)
4. P. Boersma: Praat, a system for doing phonetics by computer. Glot International 5(9/10), 341–345 (2001)
5. A.T. Corbett, J.R. Anderson: Knowledge tracing: Modeling the acquisition of procedural knowledge. User Modeling and User-Adapted Interaction 4(4), 253–278 (1994)
6. J. Schulman, F. Wolski, F. Dhariwal, A. Radford, O. Klimov: Proximal Policy Optimization Algorithms. arXiv:1707.06347 (2017)
7. R. Portelas, C. Colas, L. Weng, K. Hofmann, P.-Y. Oudeyer: Automatic Curriculum Learning For Deep RL: A Short Survey. In: Proceedings of IJCAI, pp. 4819–4825 (2020)
8. D.J. Weiss: Improving measurement quality and efficiency with adaptive testing. Applied Psychological Measurement 6(4), 473–492 (1982)
9. Z. Ji et al.: Survey of Hallucination in Natural Language Generation. ACM Computing Surveys 55(12), Article 248 (2023)
10. Qwen Team: Qwen2.5 Technical Report. arXiv:2412.15115 (2024)
11. N. Gomez, S.S. Batham, M. Volonte, T.D. Do: Virtual Interviewers, Real Results: Exploring AI-Driven Mock Technical Interviews on Student Readiness and Confidence. In: Proceedings of ACM CSCW Companion 2025, pp. 209–213 (2025). doi:10.1145/3715070.3749227
12. Q. Wu et al.: AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. arXiv:2308.08155 (2023)

13. S. Hong et al.: MetaGPT: Meta Programming for A Multi-Agent Collaborative Framework. In: International Conference on Learning Representations (ICLR) (2024)
14. E. Sun, L. Tai: MultiTutor: Collaborative LLM Agents for Multimodal Student Support. In: Proceedings of the Innovation and Responsibility in AI-Supported Education Workshop, PMLR 273, pp. 174–190 (2025)
15. M. Cemri et al.: Why Do Multi-Agent LLM Systems Fail? arXiv:2503.13657 (2025)
16. J.H. Saltzer, M.D. Schroeder: The protection of information in computer systems. Proceedings of the IEEE 63(9), 1278–1308 (1975). doi:10.1109/PROC.1975.9939
17. M. Alshiekh et al.: Safe Reinforcement Learning via Shielding. In: Proceedings of AAAI (2018)
18. C.A. Field, A.H. Welsh: Bootstrapping clustered data. Journal of the Royal Statistical Society: Series B 69(3), 369–390 (2007). doi:10.1111/j.1467-9868.2007.00593.x
19. B.S. Bloom (ed.): Taxonomy of Educational Objectives: The Classification of Educational Goals. Longmans, Green, New York (1956)
20. L.W. Anderson, D.R. Krathwohl (eds.): A Taxonomy for Learning, Teaching, and Assessing: A Revision of Bloom's Taxonomy of Educational Objectives. Longman, New York (2001)