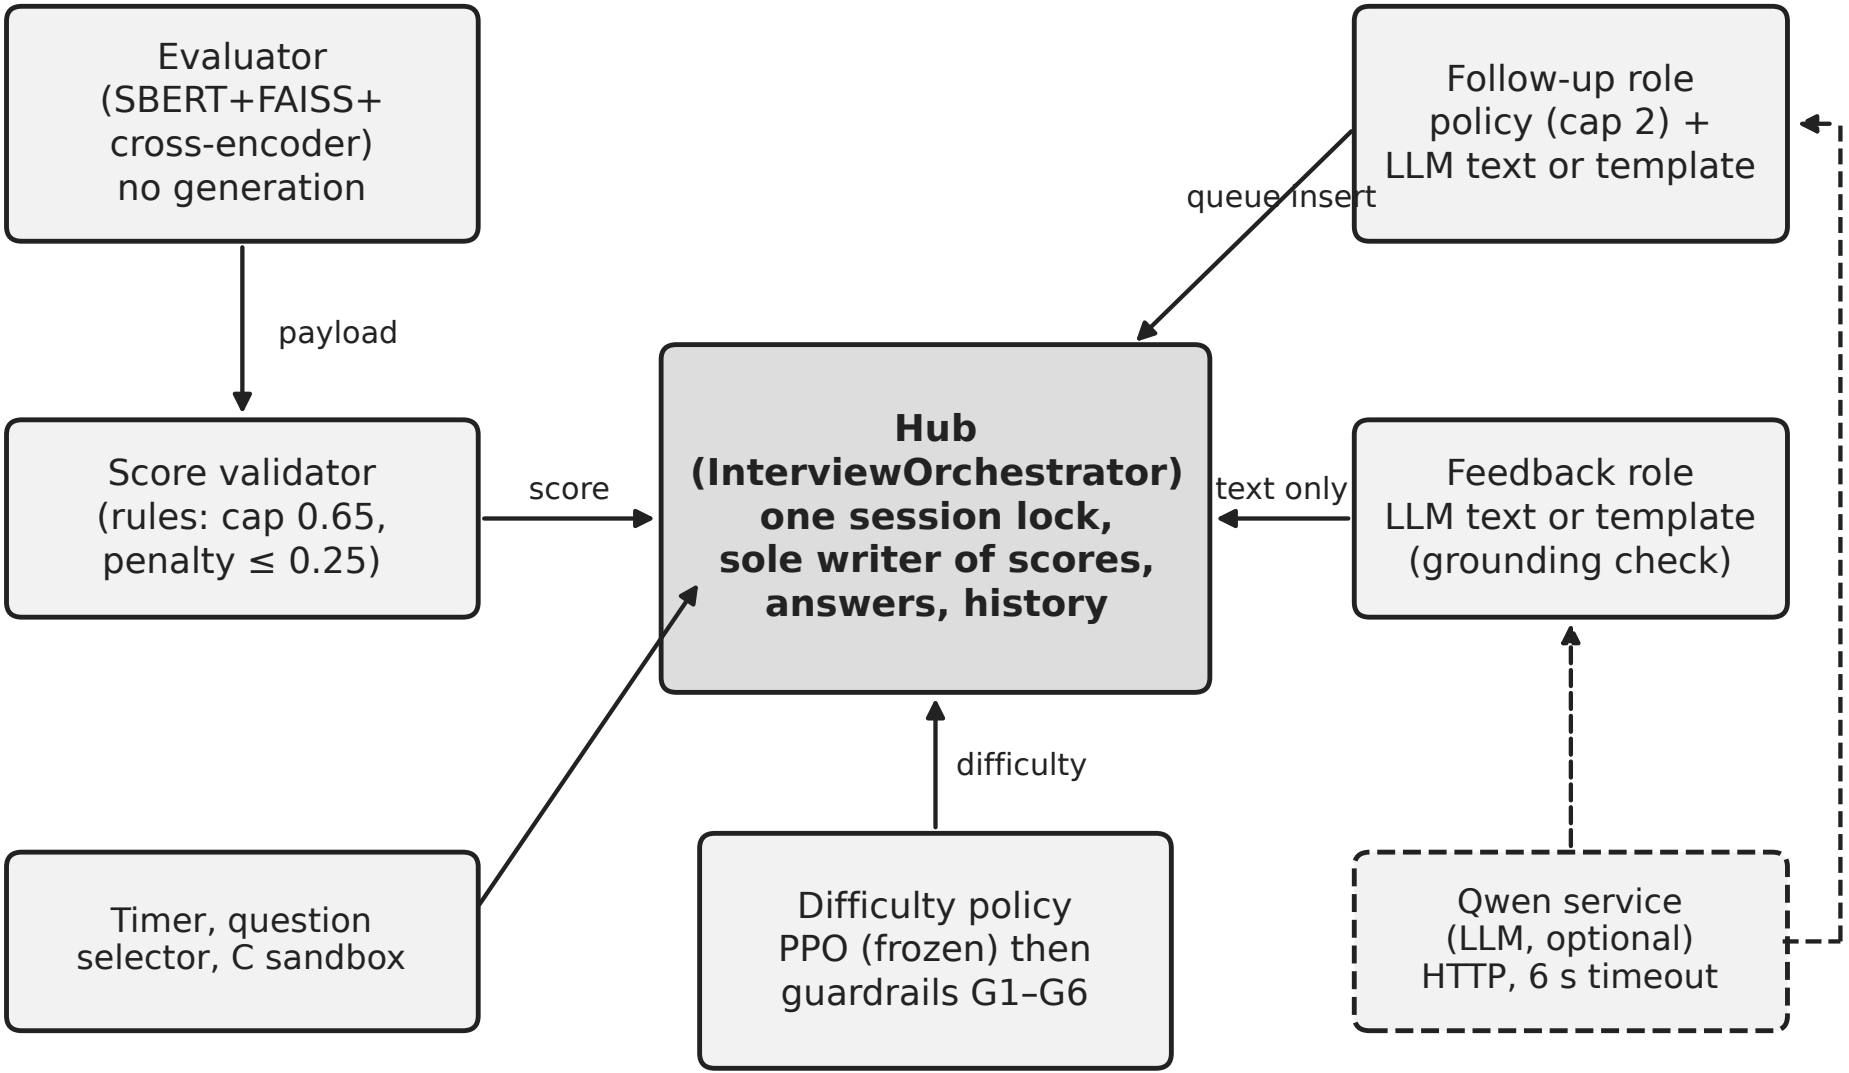
Evaluator
(SBERT+FAISS+
cross-encoder)
no generation
payload
Score validator
(rules: cap 0.65,
penalty ≤ 0.25)
score
Hub
(InterviewOrchestrator)
one session lock,
sole writer of scores,
answers, history
queue insert
Follow-up role
policy (cap 2) +
LLM text or template
text only
Feedback role
LLM text or template
(grounding check)
difficulty
Timer, question
selector, C sandbox
Difficulty policy
PPO (frozen) then
guardrails G1–G6
Qwen service
(LLM, optional)
HTTP, 6 s timeout