

Scoped Evaluation of Containment, Authority Boundaries, and Failure Handling in an LLM-Assisted Technical-Assessment System

Dr. Uma D
Professor
Dept. of CSE
PES University

Naveen S Khadd
Dept. of CSE
PES University

Sparsh Kumar
Dept. of CSE
PES University

Athreya Shashidhara
Dept. of CSE
PES University

Manasa S A
Dept. of CSE
PES University

Abstract—LLM-assisted technical-assessment pipelines execute candidate code and place model-written text near scoring, so an operator needs to know which containment, authority-boundary and failure-handling properties have been demonstrated, and by what evidence. We report a scoped empirical evaluation of one implemented pipeline (a Docker-based C sandbox, an evidence-based answer evaluator and a local 1.5B-parameter language-model feedback service) in one environment: Windows 11, WSL2 and Docker Desktop. Three campaigns used host- or runtime-side observables, five repetitions per scenario and, where possible, deliberately weakened controls. Under the shipped configuration, all nine fixed attack programs met their prespecified criteria in 5/5 runs and a benign program was accepted in 10/10; the seven attacks that had a permissive control breached in 5/5 runs under it. For five of the seven controlled attacks the executor status was identical in shipped and permissive runs, so host or runtime observables were needed. The `ptrace` result reflects a literal pre-flight filter and is not evidence of kernel-level `ptrace` containment. Two defects on the baseline system (an evaluator outage stored as a 0.0 score; a compiler timeout leaving a running container) were absent on the repaired system in 5/5 repetitions. In a fixed-turn matched-pair test, 14 measured evaluator observables were identical across 72/72 valid pairs; the follow-up-question channel was not tested. Non-numeric or negative evaluator outputs are still sanitized to 0.0, and the shipped 6 s model timeout is shorter than the measured 19–30 s generation time. These are repeatability observations in one environment, not security or fault-tolerance guarantees.

Keywords—containment testing; authority boundaries; fault injection; LLM-assisted assessment; test oracles; dependability evaluation

I. INTRODUCTION

An LLM-assisted technical-assessment pipeline does two things that deserve separate scrutiny. It compiles and runs code written by the person being assessed, and it places text produced by a language model in the same response path as a numeric score. An operator who deploys such a pipeline needs to know what has been demonstrated about three things: whether the execution environment contains what it is meant to contain, whether model-written text can write to the fields that decide scores and difficulty, and what the pipeline does when a component fails.

The pipeline studied here supports technical-interview practice: a learner submits code and written answers, and the system scores them and returns generated feedback. In such formative use, a score is only as trustworthy as the environment

that produced it. Submitted code that escapes or damages its sandbox can change the recorded result or the host, model-written feedback that reaches the score fields would let generated text rather than the answer decide the outcome, and an unflagged component failure can be recorded as a genuine zero. These are questions of assessment integrity and of the infrastructure that educational technology relies on. They are not measures of learning: the paper reports no learning, grade, usability or outcome measure.

Existing work supplies methods for each part: sandbox evaluation with fixed programs or adaptive agents judged by host-side signals [1]–[3], [6], [7]; fault injection for services and LLM-agent systems [8]–[10]; and studies of manipulation of LLM graders [11]–[13]. What these studies leave open, for one implemented pipeline, is a property-by-property statement of what was demonstrated, with which oracle, on which system and in which environment, and where the test cannot tell. Each concern here becomes a named property, a test condition, an oracle that does not rely on the component under test and, where possible, a deliberately weakened control that shows the oracle can register a failure.

We ran three campaigns: containment of nine attack programs in the code sandbox (X1-C), fault injection for failure handling (X1-A), and a fixed-turn invariance test of score-related observables against language-model-written feedback (X1-B-I). Every scenario was repeated five times as a deterministic repeatability check, not a random sample; no inferential statistics are attached.

The research question is: *Which containment, authority-boundary, and failure-handling properties of an LLM-assisted technical-assessment pipeline can be demonstrated through controlled tests under a specified execution environment?* The contributions are scoped empirical evidence of four kinds: (1) containment of nine specified attack programs under one shipped configuration, with permissive controls where a configurable control exists; (2) tested failure handling for an evaluator outage and a compiler timeout, on a baseline and a repaired system; (3) a fixed-turn observation of the language-model feedback channel with mutation and static controls; and (4) an observation that executor status strings did not identify containment for five attacks. We do not claim that the pipeline is secure, that the sandbox is universally isolating, or that model text cannot influence scoring.

II. RELATED WORK

Containers share the host kernel, and containment depends on configuration and threat model [3]. Sandbox benchmarks score escape by host flag retrieval, with deliberately introduced weaknesses [1], or by outcomes observed by test code across 51 manually written cases [2]; a comparative study separates pass, fail, partial, inconclusive and skipped verdicts [3]; commentary recommends state checks and canary tokens over an agent's own report [6], [7]; a framework proposes composing evidence into a bounded claim [4]; and a survey attributes high denylist failure rates to a third-party study [5]. Injecting named faults is established for services [8] and has been applied to LLM-based agent systems [9], [10], whose evaluations concern task success; separating model output from privileged actions is an established pattern [11]; and evaluations of LLM graders report that instructions injected into a submission can change the grade [12], [13]. The present study applies fixed programs, host-side observables, weakened controls and bounded verdicts to one shipped configuration and to one model-text channel. It does not measure escape capability against adaptive attackers.

III. PIPELINE, SCOPE AND THREAT MODEL

The pipeline comprises an orchestrator holding session state, question queue, score and difficulty fields; an evaluator service producing the numeric score; a local 1.5B-parameter instruction-tuned language model (quantized GGUF, CPU) writing narrative feedback and follow-up questions; a Docker-based C compiler-and-runner sandbox; and a SQLite store. Three authority boundaries are examined: (a) code from the assessed person runs only inside the sandbox; (b) score, difficulty and best-answer fields are written by the evaluator path, not by model text; (c) a failed component surfaces as a flagged, unscored outcome rather than a fabricated score. These are design intentions; Section V reports what the tests could show about each. The follow-up-question generator is a fourth path, because its output can become the expected-concept list of the next turn; it was not tested.

Threat model and scope. The tested adversary is a candidate who submits a fixed C program attempting one of nine behaviours (Table II), or a benign or injection-carrying transcript for the feedback model. Adaptive attackers, kernel or runtime exploits, UDP/DNS/IPv6 paths and other platforms are outside scope. "Containment" here means that the recorded outcome of a named program under the shipped configuration met prespecified criteria. Some results are produced by container settings (network, memory, process limit, read-only filesystem) and one by an application-level pre-flight filter; each is attributed to its layer below.

The shipped configuration, read from the recorded Docker commands, is: user 1001:1001, --net=none, all capabilities dropped, no-new-privileges, read-only root filesystem, tmpfs workspace, 128 MB memory and swap, 32 processes, 1 CPU, program timeout 2 s, and Docker's default seccomp profile.

IV. METHOD

Systems under test. The *baseline SUT* is the tag `release/app-repair/v1` (commit 980747ff). The *repaired SUT* is `sut/X1/build-B` (commit beb374f3), which repairs the evaluator-outage and compiler-timeout defects found on the

baseline, makes the model-status flag truthful and caps feedback generation at 512 tokens. (The similarly named tag `sut/X1/build-A` belongs to a different, unreported campaign and is not the baseline.) Both systems are reported.

Design and oracles. Attack programs, controls and oracles were written by the study's authors. Table I lists property, test, oracle and control. An oracle combines a program's self-report with host-side observables the program cannot forge: a canary listener on the host (connection count), a hash of a canary directory, `docker events` (out-of-memory and kill events) and an orphan-container check.

TABLE I. PROPERTIES, TESTS, ORACLES AND CONTROLS.

Property	Test	Oracle	Control
Containment of nine fixed attack programs	X1-C: each program under the shipped configuration	Conjunction of self-report and host-side observables	Permissive configuration (one documented change) for seven of nine; benign program
Feedback text does not change score, difficulty or best-answer fields	X1-B-I: 36 injections $\times$ 2 replicates, benign and injection arms, evaluator output fixed	Exact equality of 14 stored observables (scores, grade, breakdown, difficulty, next action, best-answer flags and the stored best-answer score field, queue length)	Mutation control (a harness copy wiring a model-text field into the score); static AST guard
Evaluator outage and compiler timeout are handled without a fabricated score or orphan container	X1-A FLT-03, FLT-06 (plus seven other faults)	Infrastructure flag present and no 0.0 in session scores; no container 3 s after return	No-fault control

Pre-specification and status. Protocols were committed and tagged with local annotated git tags before execution; the tags were not pushed, so this pre-specification is author-controlled and not externally time-stamped. Reruns on the repaired SUT (X1-C v2, X1-A v2, X1-B-I v3) were written after the baseline results were seen. This is a scoped engineering evaluation, not a preregistered inferential study. "k/5" counts runs meeting a criterion and is not a rate.

Environment. Windows 11 (10.0.26200), WSL2 kernel 6.6.87.2, Docker Desktop with Docker 29.7.2, runc, cgroup v2, 12 CPUs, 8.15 GB memory; the project's local sandbox image; language model `qwen2.5-1.5b-instruct-q4_k_m` with default decoding (temperature 0.1, top-p 0.9) and no fixed seed. No other environment was tested.

Fixed-turn matched-pair test. The evaluator output is replaced by a fixed record (score 0.90), so a change in the compared observables can only come from the feedback path. Each pair runs the same answer twice, with a benign or an injection-carrying transcript; a pair is valid only if both arms' feedback was produced by the model rather than a template. The pre-specified minimum was 30 valid pairs. The client timeout was raised to 600 s for the test (a disclosed deviation from the shipped 6 s).

V. RESULTS

A. Containment (X1-C)

Fig. 1 shows the per-run outcomes; Table II gives what was observed and which observable decided each outcome. On both

systems all nine programs met the criteria under the shipped configuration in 5/5 runs; the seven permissive controls breached in 5/5 runs; the benign program was accepted in 10/10 runs (five shipped, five permissive-all).

Attack program	Shipped (criteria met)	Permissive control (breach detected)	Executor status same in both?	Status string(s)
SEC-01 ptrace	●●●●●	*****	differs	policy_blocked wrong_answer
SEC-02 socket + connect	●●●●●	*****	identical	wrong_answer wrong_answer
SEC-03 syntax error	●●●●●	no control	-	compilation_error
SEC-04 SIGSEGV	●●●●●	no control	-	runtime_error
SEC-05 CPU loop	●●●●●	*****	identical	timeout timeout
SEC-06 memory	●●●●●	*****	differs	memory_limit wrong_answer
SEC-07 exfiltration	●●●●●	*****	identical	wrong_answer wrong_answer
SEC-08 fork attempts	●●●●●	*****	identical	wrong_answer wrong_answer
SEC-09 filesystem write	●●●●●	*****	identical	wrong_answer wrong_answer
BEN benign program	●●●●●●●●●● (in shipped cell)	-	-	accepted

● shipped run met all criteria * permissive run breached Repaired SUT, 5 runs per cell; baseline SUT identical.

Fig. 1. Per-run X1-C outcomes on the repaired SUT. Blue dot: shipped run met all prespecified criteria; vermilion cross: permissive-control run breached. Executor status strings were identical in shipped and permissive runs for SEC-02, -05, -07, -08 and -09. Source: runs.jsonl, results.csv of X1-C v2; the baseline SUT gave identical counts.

TABLE II. SHIPPED OBSERVATION AND DECIDING OBSERVABLE (REPAIRED SUT; BASELINE IDENTICAL).

Attack	Shipped observation	Deciding observable / layer
SEC-01 ptrace	rejected before execution (policy_blocked)	literal pre-flight filter (application layer); see below
SEC-02 socket + connect	socket() = 3; connect() = -1 (errno 101)	host canary listener: 0 connections; --net=none removed in the control
SEC-03 syntax error; SEC-04 SIGSEGV	compilation error; exit 139	executor status (no permissive control)
SEC-05 CPU loop	timeout, exit 124; execution time 2.63–2.89 s (baseline 2.53–2.56 s), bound 6 s	execution time; control ran to its extended 12 s limit
SEC-06 memory	killed, exit 137	docker events out-of-memory event
SEC-07 exfiltration	connect failed, 0 bytes	canary listener: control received the 39-byte token
SEC-08 fork attempts (≤100)	30 forks, then EAGAIN	fork count vs. configured limit; control reached 100
SEC-09 filesystem write	write under /etc failed (EROFS)	write errno; canary-directory hash unchanged (see caveat)

ptrace. SEC-01 tested only the literal `ptrace()` spelling. The shipped pre-flight filter rejected the source text. With the filter disabled, a direct `ptrace(PTRACE_TRACEME)` call returned 0 inside the container. Docker documents `ptrace` as blocked in kernels before 4.8 and states that tracing arbitrary processes is already blocked by dropping `CAP_SYS_PTRACE` [14]; the tested kernel is 6.6.87.2. The observation is compatible with that documentation but shows that, in the tested path, containment of this attack depended on the filter. It is therefore not evidence of kernel-level

`ptrace` containment, and this paper does not attribute the SEC-01 outcome to Docker.

Oracle observations. The executor reported `wrong_answer` in both the shipped and permissive configuration for SEC-02, SEC-07, SEC-08 and SEC-09, so contained and breached runs could not be told apart from the status string; the observables in Table II decided them. For SEC-05 the status was `timeout` in both configurations and execution time separated them, so the status did not separate the configurations for five attacks in all. `socket()` succeeded and `connect()` failed, so "socket creation blocked" would misdescribe SEC-02. The SEC-08 oracle references the configured process limit and is therefore not fully independent of the system under test. The SEC-09 canary criterion does not discriminate in the shipped configuration and the permissive control's canary reset was flawed, so only the write-failure observation is relied on. Execution times are measurements on this machine and are not latency claims.

B. Failure handling (X1-A)

TABLE III. X1-A CRITERIA MET OF FIVE (BASELINE SUT / REPAIRED SUT).

Scenario	Baseline	Repaired	Note
FLT-03 evaluator outage	0/5	5/5	baseline: stored as 0.0, no flag; repaired: evaluator_unavailable, unscored
FLT-06 compiler outlives timeout	0/5	5/5	baseline: container present 3 s after return; repaired: none
FLT-07b database lock 34 s	4/5	5/5	one baseline run 0.1 s above the bound; not a repair
FLT-01, -02, -04, -05, -07a, -10	5/5 each	5/5 each	FLT-04: NaN, Inf and -3 stored as unflagged 0.0 (unrepaired); 999 → 1.0
FLT-08 WebSocket; FLT-09 audio	not executed	not executed	no result

In Table III, the two defects are the substantive findings. On the baseline SUT the injected evaluator outage was stored as a 0.0 score with no flag, indistinguishable in stored state from a genuine zero; on the repaired SUT it produced an infrastructure flag and no score, answer, attempt or difficulty change (5/5). The baseline left a running container after a compiler timeout in 5/5 runs; the repaired SUT left none. Both repairs were designed, applied and re-tested by the same AI coding agent that found the defects, so the re-test is a regression check and not an independent replication. FLT-04 shows only that scores ended inside [0, 1]: NaN, Inf and negative outputs become 0.0, the outcome FLT-03 identified as a defect when caused by an outage, so FLT-04 is a range-sanitization observation and not evidence of corruption detection. FLT-08 and FLT-09 were not executed. FLT-01 and FLT-02 used a local model stub and do not test the real service.

C. Language-model text and evaluator observables (X1-B-I, fixed turn)

TABLE IV. FIXED-TURN MATCHED-PAIR TEST.

Quantity	Baseline SUT	Repaired SUT
Pairs / valid pairs / pre-specified minimum	72 / 16 / 30 (not met)	72 / 72 / 30 (met)
14 observables exactly equal (valid pairs)	16/16	72/72
Mutation control detected	50/72	72/72
Static guard: real sources flagged / injected mutant caught	0 of 75 / yes	0 / yes
Arms served by template instead of the model	78 of 144	0 of 144

Table IV summarises the test. Measured evaluator observables were unchanged across the 72 valid matched pairs on the repaired SUT under the tested fixed-turn perturbation, the model text differed between arms in 72/72 pairs, and the mutation control detected the injected authority path in 72/72. **The downstream follow-up-question generation channel was not tested.** On the baseline SUT only 16 of 72 pairs were valid because template-served feedback was reported as model-produced, the pre-specified minimum was not met, and that result is diagnostic only. The steering effect of the injections was not demonstrated: a crude check found an injection-specific word in the adversarial-arm text in 23 of 72 repaired-SUT pairs, and a test in which the model did not follow the injections would

pass trivially. The test covers narrative feedback with the evaluator output held fixed.

D. Timing of the model service

The shipped model-client timeout is 6 s; local CPU generation took about 19–30 s per valid benign arm in X1-B-I, which used a 600 s timeout. These times do not show that the shipped 6 s path completes generation, and no end-user consequence was tested.

VI. DISCUSSION

Containment. Under the tested configuration, harness and machine, each of the nine fixed programs met its criteria, and the permissive controls show that the oracles can register a breach. The result says nothing about adaptive attackers, kernel or runtime escapes, other network paths or platforms, or programs outside the nine. Host-side observables, weakened controls and repetition follow benchmark practice [1], [6], [7]; nothing here is a new escape result. A passing status would not have told an operator that SEC-01 rested on a literal filter, so a pass for a named program should be attributed to the layer that produced it.

Executor status as oracle. In this harness the status string was identical across contained and breached runs for five attacks (SEC-02, -05, -07, -08 and -09; for the CPU loop only the execution time differed), consistent with the advice to check state rather than trust reported outcomes [6], [7]; this does not show that status strings are uninformative in general.

Failure handling. Two specific defects existed on the baseline SUT and were absent on the repaired SUT in the tested scenarios; an unflagged 0.0 for an outage is the kind of failure a scoring pipeline should surface. This does not show general fault tolerance, and untested faults, hangs longer than tested and the unflagged sanitization of NaN, Inf and negative scores remain.

Model authority. Grader-injection studies [11]–[13] concern models that grade; here the model does not grade, and the question is whether its text reaches score observables. The evidence covers one channel, with fixed evaluator output, 36 single-sentence injections and unseeded generation.

The study reports observed behaviour of specified scenarios on named systems; it is closer to dependability testing of a pipeline with security-relevant components than to a security evaluation, makes no adversary-capability claim and carries no proof.

VII. THREATS TO VALIDITY AND LIMITATIONS

Four material limitations. (1) *Numeric sanitization*: NaN, Inf and negative evaluator outputs become 0.0, not infrastructure failures. (2) *Untested follow-up channel*: model-generated follow-up text can enter the evaluator’s expected-concept path and X1-B-I did not test it. (3) *ptrace*: SEC-01 depended on literal pre-flight filtering; the bypassed call returned 0 in the container. (4) *Timeout*: the shipped 6 s client timeout versus 19–30 s local generation.

Construct and internal validity. Attack programs, oracles and controls are author-written; the pre-flight is a literal-pattern filter tested with one spelling; the SEC-08 oracle and some X1-A timing bounds derive from constants in the system under test;

the SEC-09 control had a flawed canary reset. One AI coding agent designed, pre-specified, ran, repaired and re-tested the campaigns under the direction of the project lead, and reruns on the repaired SUT were written after baseline results were seen. The X1-B-I test used a fixed evaluator stub, a name-based static guard, replayed-text mutation control, unseeded generation and a raised client timeout.

External and conclusion validity. One developer-workstation environment; nine finite programs; network tests covered one TCP connection to a host listener; enumerated faults only; FLT-08 and FLT-09 not executed; an unreachable Docker daemon was not campaign-tested. Five deterministic repetitions give repeatability counts, not rates. The development environment's package versions violate its declared pins.

Not evaluated. Real users, learning outcomes, latency guarantees, cross-platform behaviour, adaptive attackers and any property beyond the tested scenarios.

VIII. REPRODUCIBILITY AND AVAILABILITY

Harness scripts, protocols, per-run records, output hashes and the environment record are stored in the project repository, which is not publicly released in this draft; output directories are write-once and each harness checks the tagged blobs against the SUT tag. Timings are machine-dependent and model generation is unseeded, so only counts are reproducible. Artifact availability is not independent reproduction, which has not been performed.

IX. CONCLUSION

Under one Windows 11, WSL2 and Docker Desktop environment, nine fixed attack programs met prespecified criteria in 5/5 runs under the shipped configuration, and the seven permissive controls breached in 5/5 runs. Executor status alone did not identify containment for five attacks. The ptrace result reflects a literal pre-flight filter. Two failure-handling defects on the baseline system were absent on the repaired system in the two tested scenarios. Under a fixed-turn test, 14 measured evaluator observables were unchanged across 72/72

valid matched pairs; the follow-up channel and the shipped timeout were not shown to be covered. These are scoped observations, not guarantees.

REFERENCES

- [1] R. Marchand *et al.*, "Quantifying frontier LLM capabilities for container sandbox escape," arXiv:2603.02277, 2026 (preprint; the arXiv header states ICML 2026, proceedings listing not verified).
- [2] R. Rabin, J. Hostetler, S. McGregor, B. Weir, and N. Judd, "SandboxEval: Towards securing test environment for untrusted code," arXiv:2504.00018, 2025 (preprint).
- [3] G. Andronchik and P. Lokhmakov, "AI code sandboxes: A comparative security study. Part 1 of 2," arXiv:2606.08433, 2026 (preprint).
- [4] I. Singh, H. Mahmoud, and A. Murillo, "AI sandboxes: A threat model, taxonomy, and measurement framework," arXiv:2606.18532, 2026 (preprint).
- [5] M. Rashidi, "The balkanization of execution-security research for AI coding agents," arXiv:2607.05743, 2026 (preprint; abstract-level reading).
- [6] C. Guo *et al.*, "RedCode: Risky code execution and generation benchmark for code agents," arXiv:2411.07781, 2024 (preprint; venue not verified).
- [7] S. Abdelnabi, C. Hicks, K. Rieck, and A.-R. Sadeghi, "Measuring security without fooling ourselves: Why benchmarking agents is hard," arXiv:2605.22568, 2026 (preprint).
- [8] A. Basiri *et al.*, "Chaos engineering," *IEEE Softw.*, vol. 33, no. 3, pp. 35–41, 2016, doi: 10.1109/MS.2016.60.
- [9] J. Jia, Z. Deng, Z. Chen, Y. Wang, and Z. Zheng, "MAS-FIRE: Fault injection and reliability evaluation for LLM-based multi-agent systems," arXiv:2602.19843, 2026 (preprint).
- [10] A. Gupta, "ReliabilityBench: Evaluating LLM agent reliability under production-like stress conditions," arXiv:2601.06112, 2026 (preprint).
- [11] E. Debenedetti *et al.*, "Defeating prompt injections by design," arXiv:2503.18813, 2025 (preprint; abstract-level reading).
- [12] H. Li *et al.*, "'Important You should give me full credits!': Exploring prompt injection attacks on LLM-based automatic grading systems," arXiv:2606.03090, 2026 (preprint).
- [13] D. Sahoo *et al.*, "The compliance paradox: Semantic-instruction decoupling in automated academic code evaluation," arXiv:2601.21360, 2026 (preprint).
- [14] Docker Inc., "Seccomp security profiles for Docker," Docker Docs. [Online]. Available: <https://docs.docker.com/engine/security/seccomp/> (accessed Sep. 21, 2026).